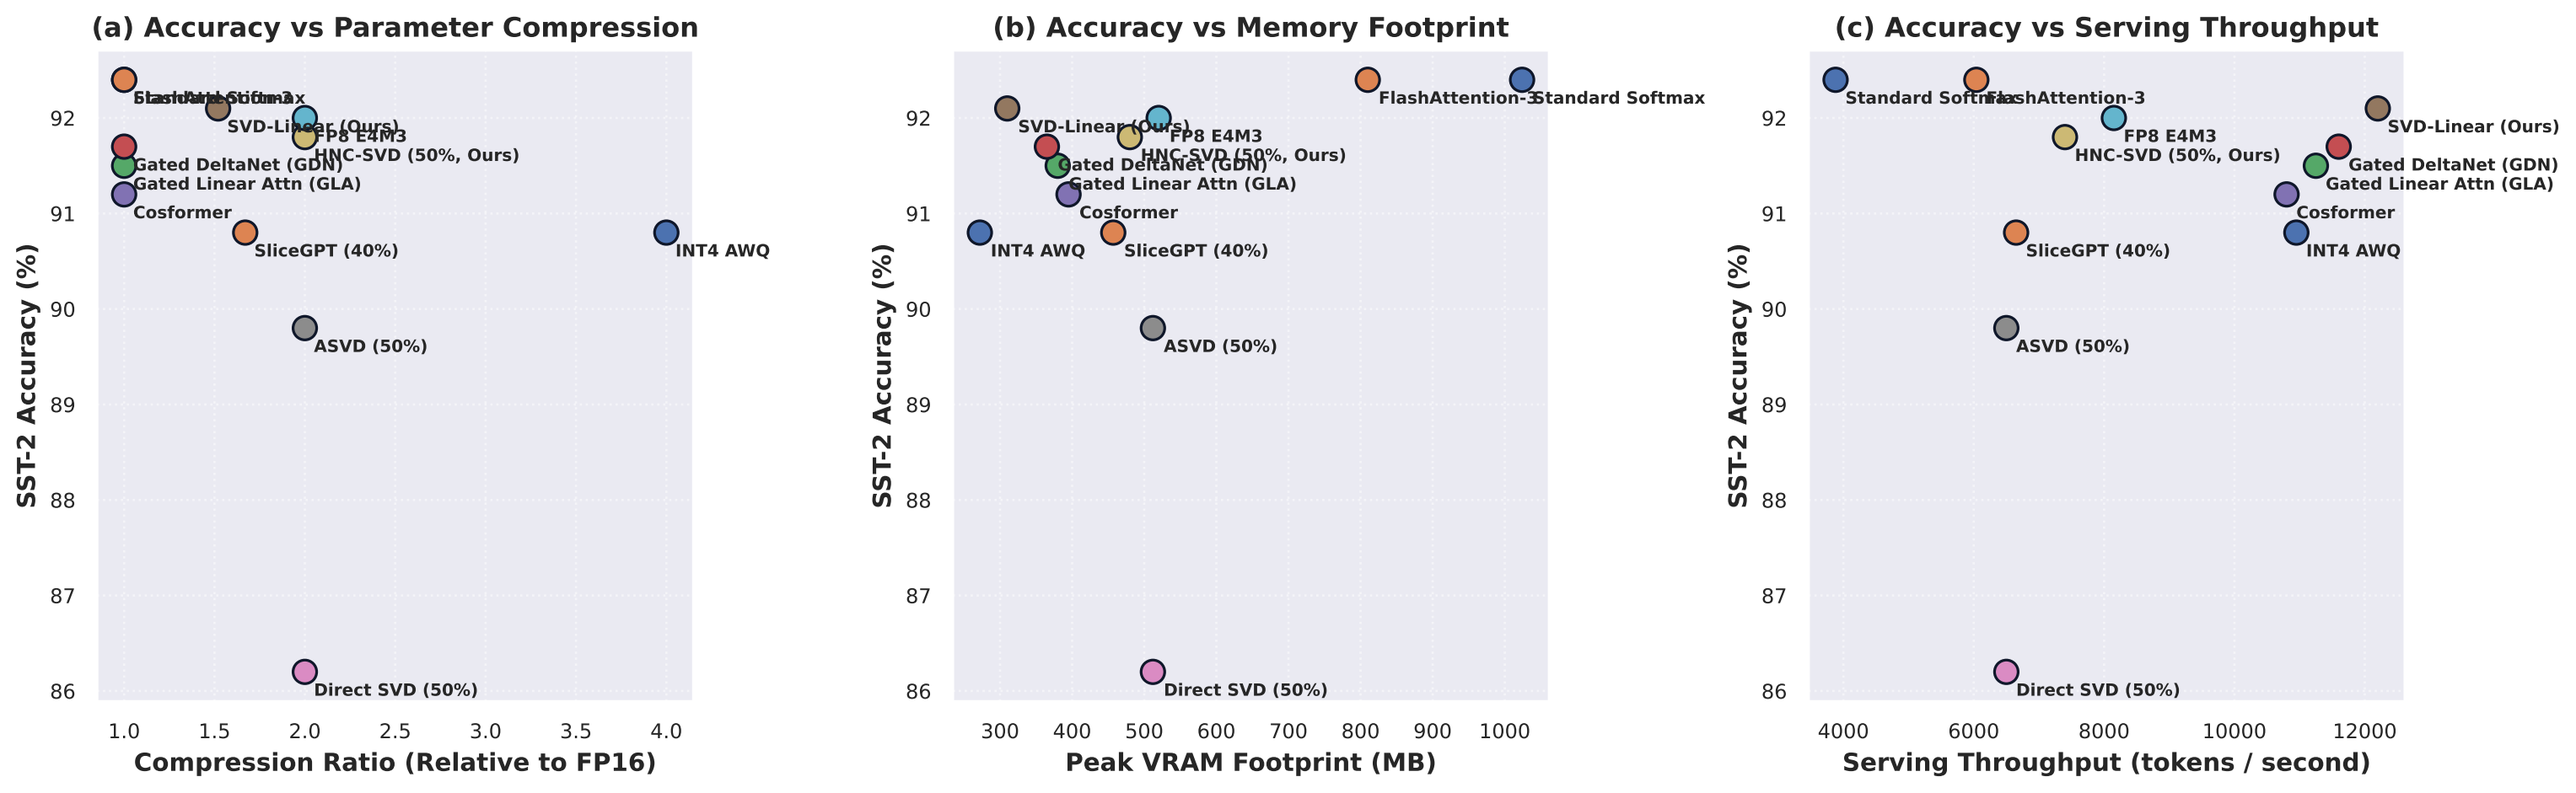
(a) Accuracy vs Parameter Compression
SST-2 Accuracy (%)
Compression Ratio (Relative to FP16)
FlashAttention-3
Standard Softmax
SVD-Linear (Ours)
FP8 E4M3
HNC-SVD (50%, Ours)
Gated DeltaNet (GDN)
Gated Linear Attn (GLA)
Cosformer
SliceGPT (40%)
INT4 AWQ
ASVD (50%)
Direct SVD (50%)
(b) Accuracy vs Memory Footprint
SST-2 Accuracy (%)
Peak VRAM Footprint (MB)
FlashAttention-3
Standard Softmax
SVD-Linear (Ours)
FP8 E4M3
HNC-SVD (50%, Ours)
Gated DeltaNet (GDN)
Gated Linear Attn (GLA)
Cosformer
INT4 AWQ
SliceGPT (40%)
ASVD (50%)
Direct SVD (50%)
(c) Accuracy vs Serving Throughput
SST-2 Accuracy (%)
Serving Throughput (tokens / second)
Standard Softmax
FlashAttention-3
SVD-Linear (Ours)
FP8 E4M3
HNC-SVD (50%, Ours)
Gated DeltaNet (GDN)
Gated Linear Attn (GLA)
Cosformer
SliceGPT (40%)
INT4 AWQ
ASVD (50%)
Direct SVD (50%)