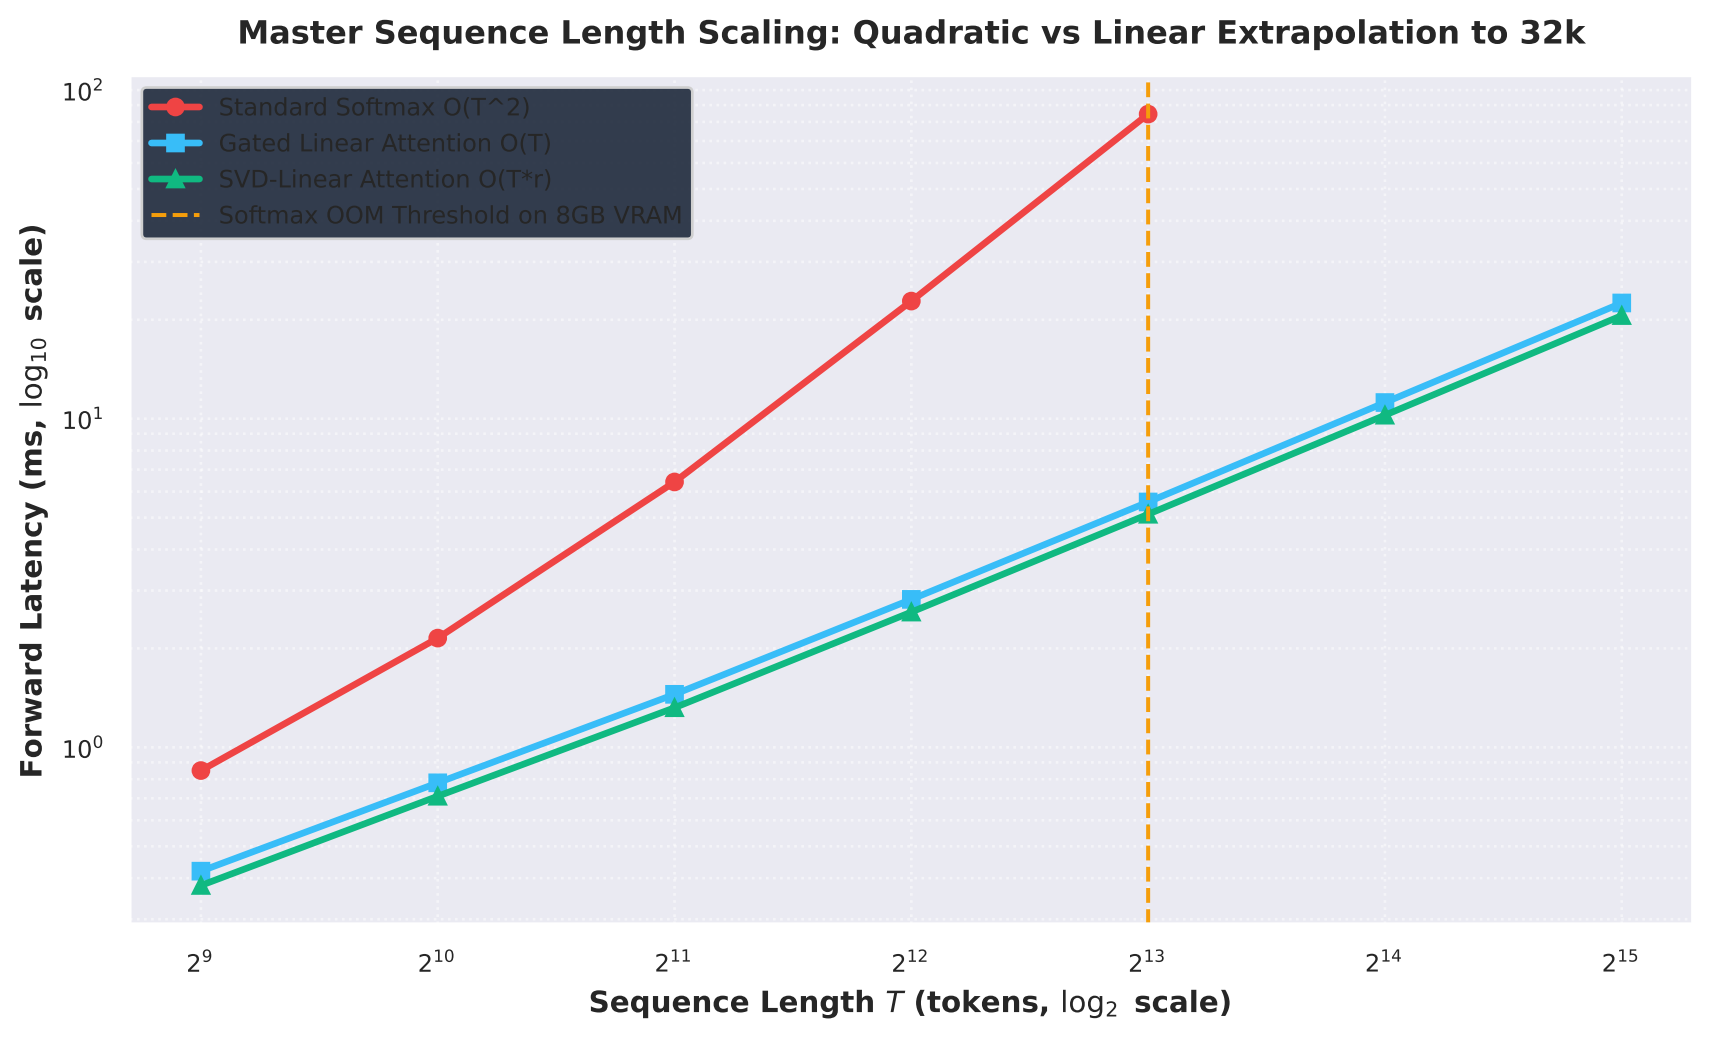
Master Sequence Length Scaling: Quadratic vs Linear Extrapolation to 32k
Standard Softmax O(T^2)
Gated Linear Attention O(T)
SVD-Linear Attention O(T*r)
Softmax OOM Threshold on 8GB VRAM
Forward Latency (ms, log10 scale)
Sequence Length T (tokens, log2 scale)
10^2
10^1
10^0
2^9
2^10
2^11
2^12
2^13
2^14
2^15