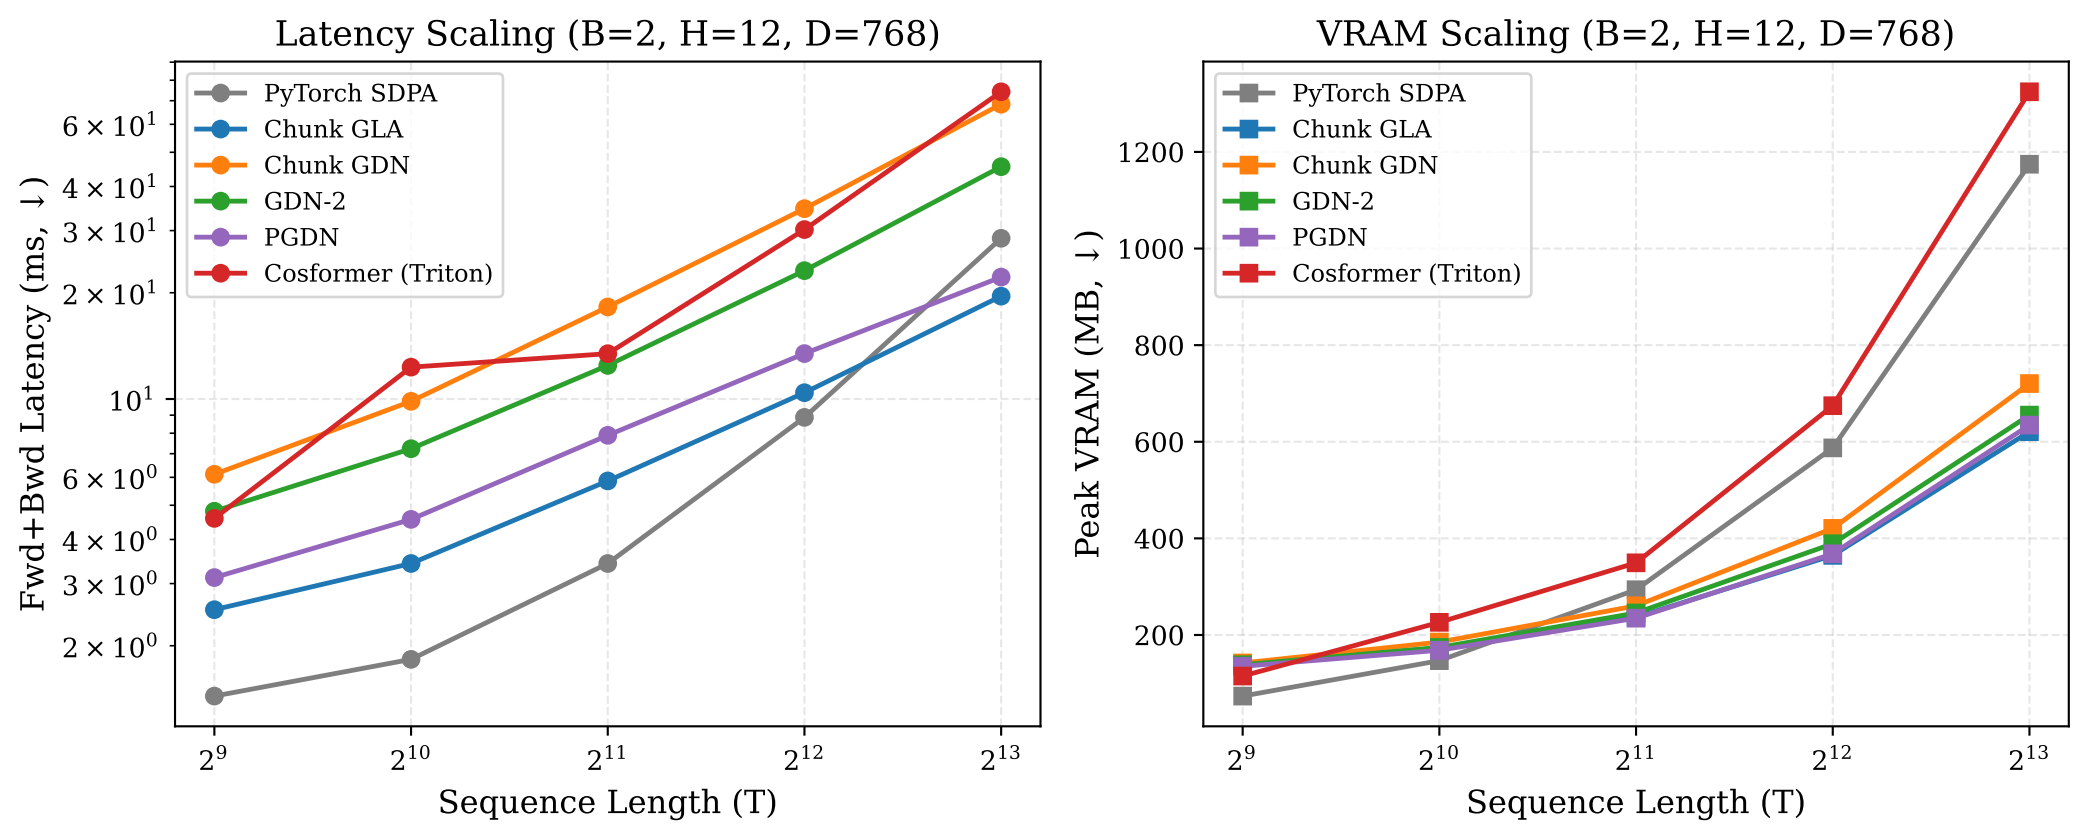
Latency Scaling (B=2, H=12, D=768)
Fwd+Bwd Latency (ms, ↓)
Sequence Length (T)
PyTorch SDPA
Chunk GLA
Chunk GDN
GDN-2
PGDN
Cosformer (Triton)
VRAM Scaling (B=2, H=12, D=768)
Peak VRAM (MB, ↓)
Sequence Length (T)