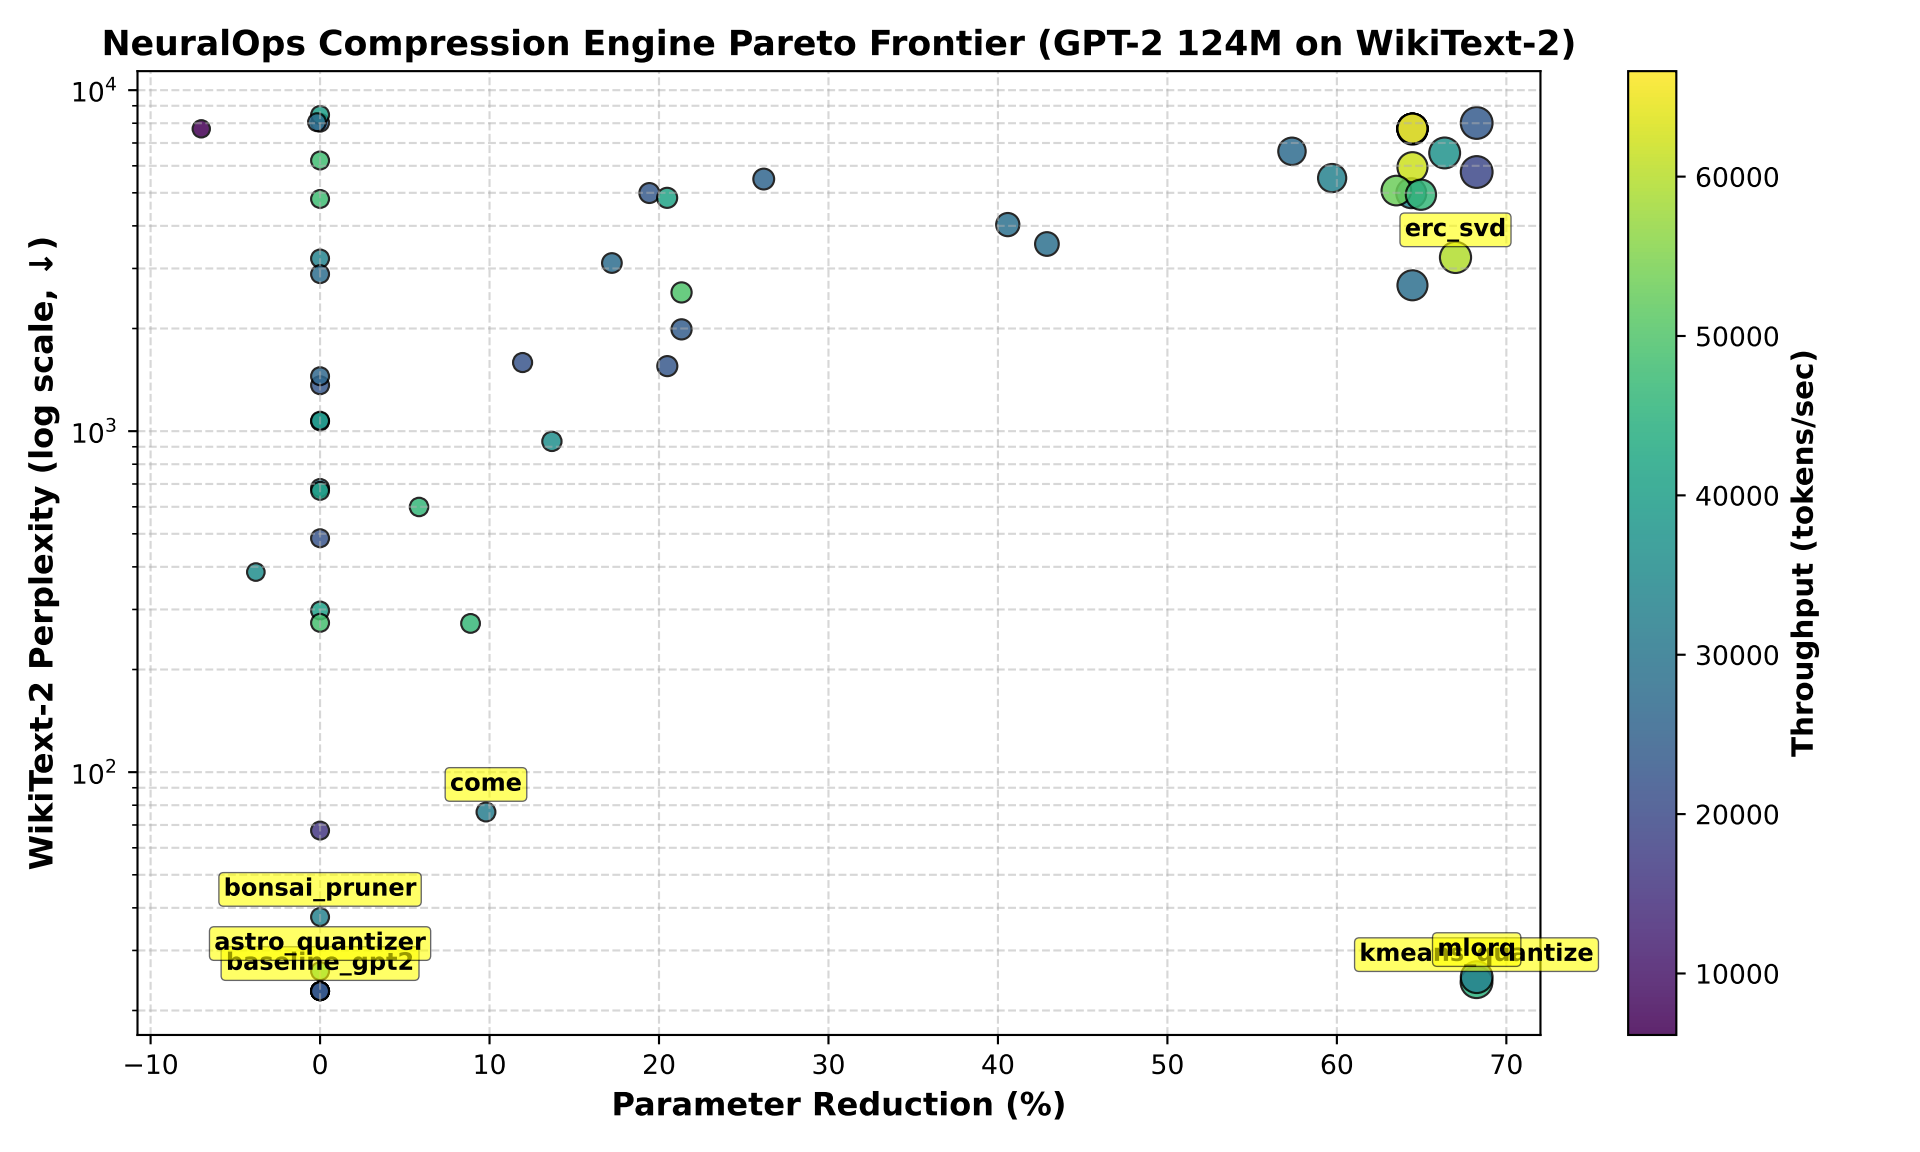
NeuralOps Compression Engine Pareto Frontier (GPT-2 124M on WikiText-2)
WikiText-2 Perplexity (log scale, ↓)
Parameter Reduction (%)
Throughput (tokens/sec)
erc_svd
come
bonsai_pruner
astro_quantizer
baseline_gpt2
mlorq