NeuralOps Evaluation Pipeline: End-to-End System Architecture

Input: Pretrained BERT / RoBERTa / T5 → Output: Per-module FP32 & FP16 Efficiency + GLUE Quality Metrics

Reproducibility: Seeds {42,123,777} | Docker pytorch:2.5.0-cuda12.1 | NVIDIA GeForce RTX 4070 Laptop GPU (8GB GDDR6, 140W TGP)