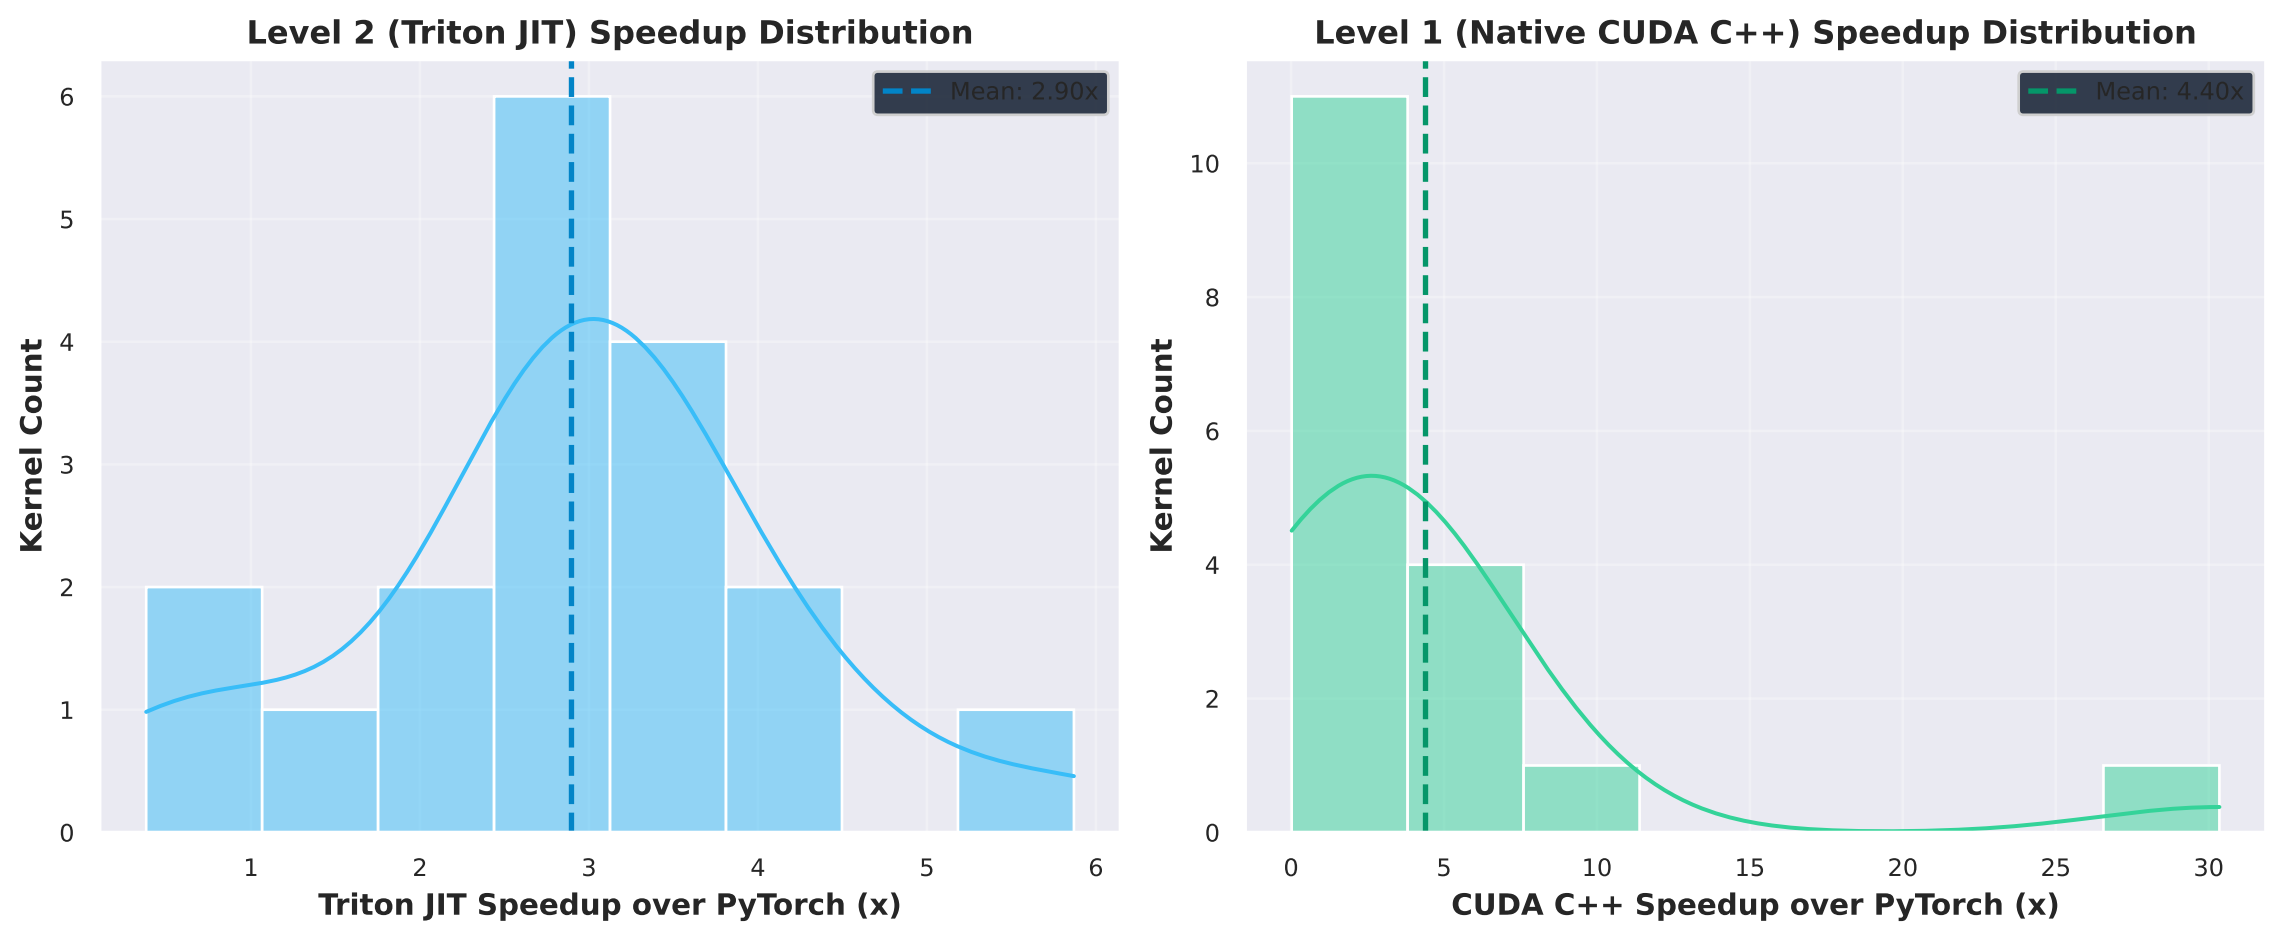
Level 2 (Triton JIT) Speedup Distribution
Mean: 2.90x
Kernel Count
Triton JIT Speedup over PyTorch (x)
Level 1 (Native CUDA C++) Speedup Distribution
Mean: 4.40x
Kernel Count
CUDA C++ Speedup over PyTorch (x)