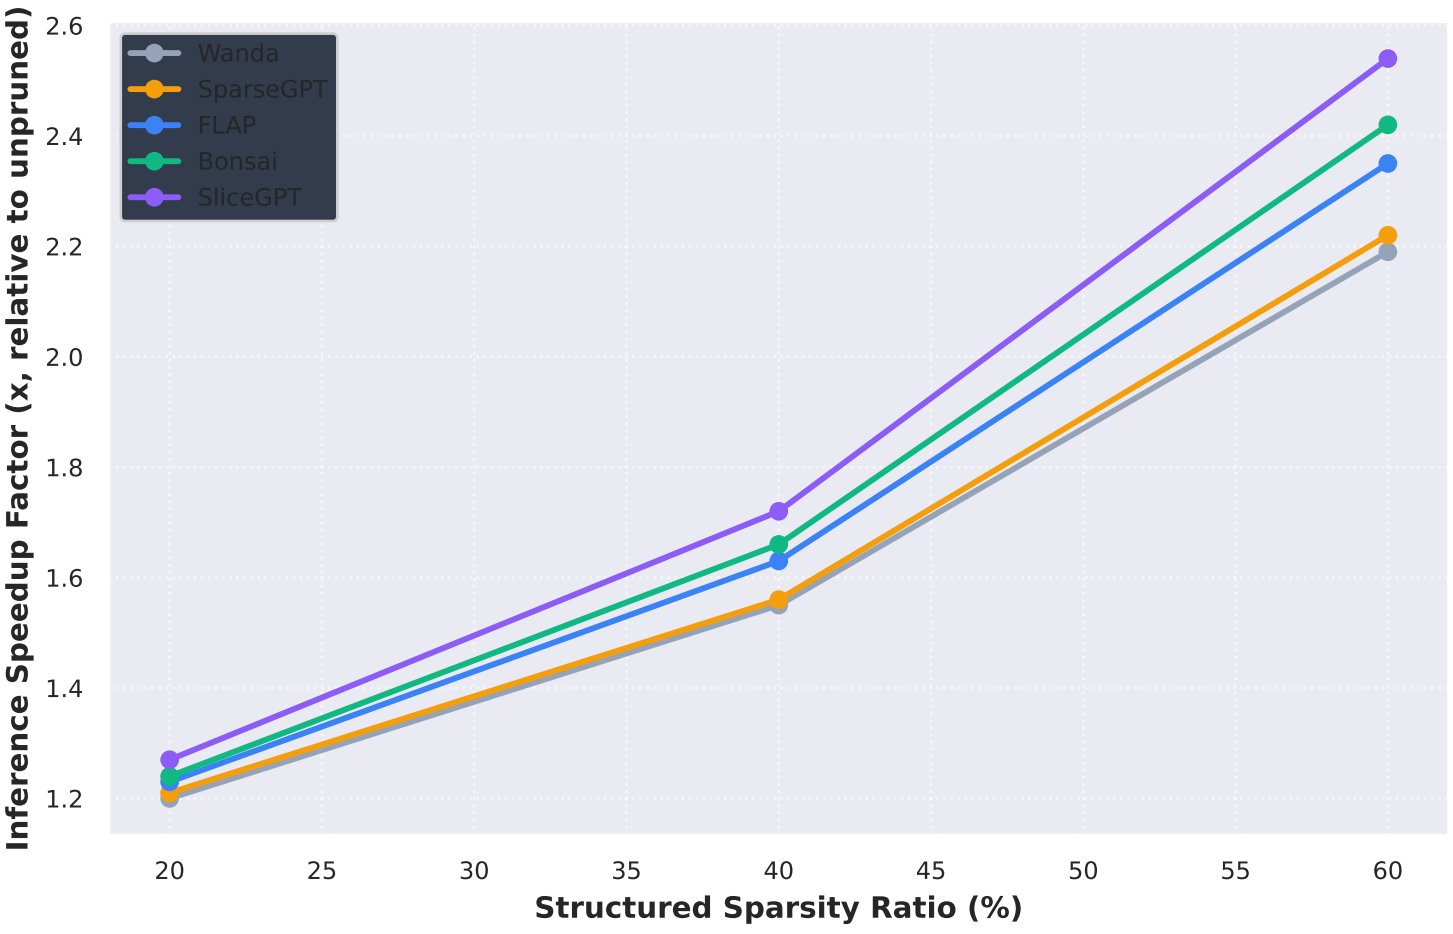
Empirical Hardware Speedup on RTX 4070 vs Pruning Sparsity
Wanda
SparseGPT
FLAP
Bonsai
SliceGPT
Inference Speedup Factor (x, relative to unpruned)
2.6
2.4
2.2
2.0
1.8
1.6
1.4
1.2
20
25
30
35
40
45
50
55
60
Structured Sparsity Ratio (%)