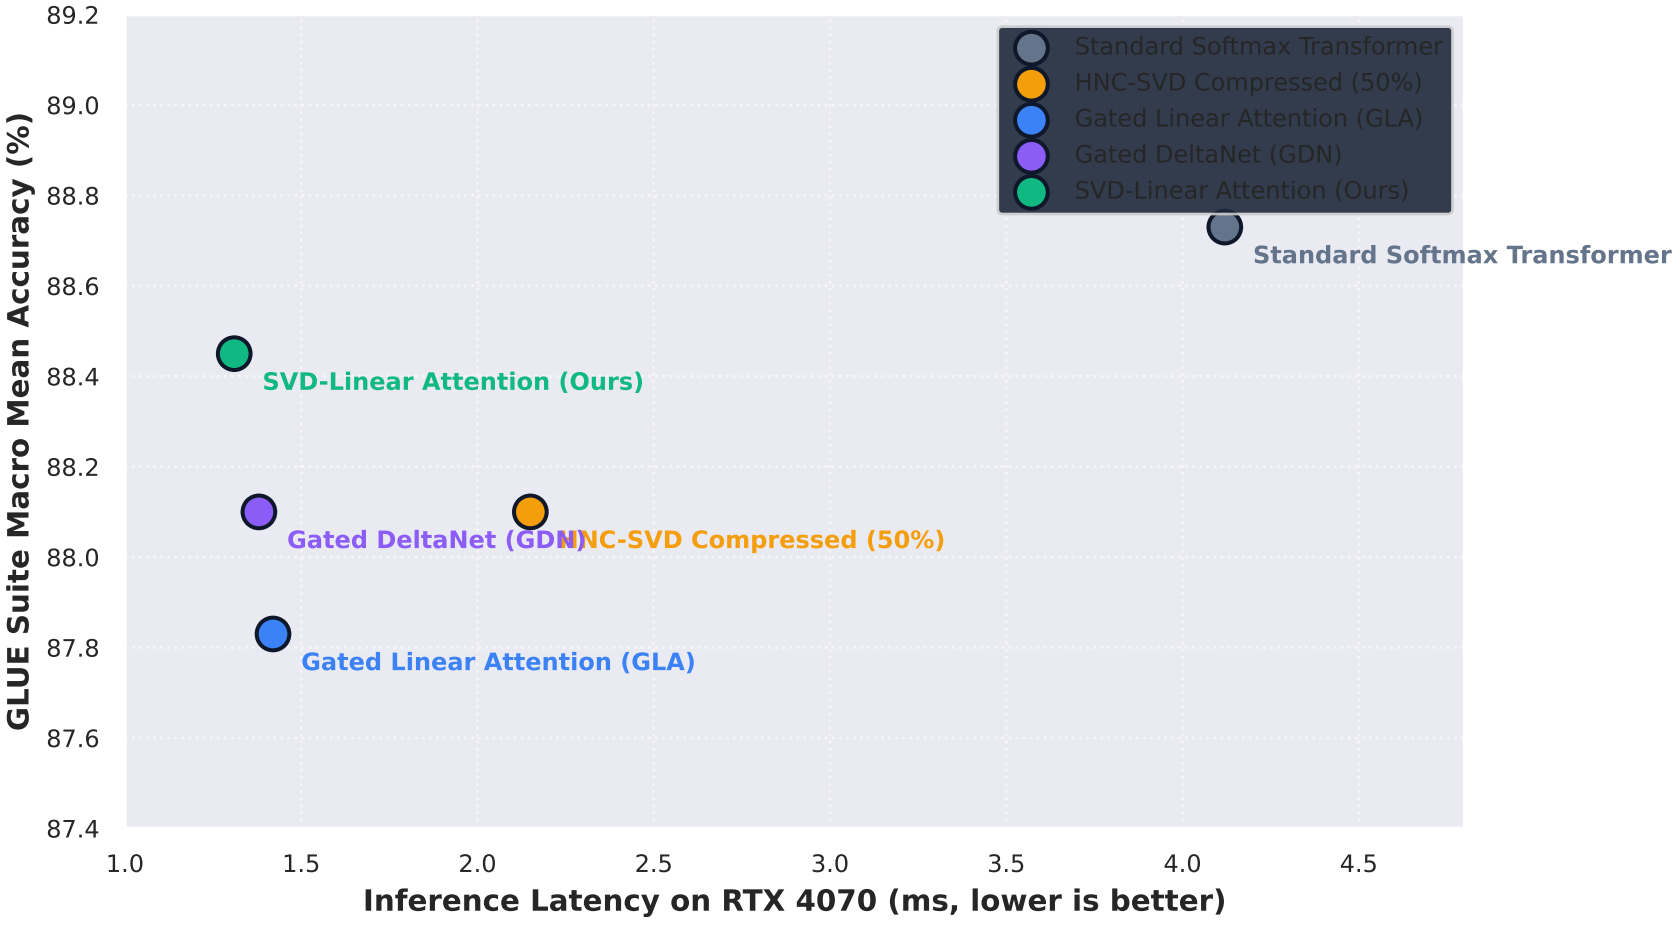
Downstream Accuracy vs Hardware Latency Pareto Frontier
GLUE Suite Macro Mean Accuracy (%)
Inference Latency on RTX 4070 (ms, lower is better)
89.2
89.0
88.8
88.6
88.4
88.2
88.0
87.8
87.6
87.4
1.0
1.5
2.0
2.5
3.0
3.5
4.0
4.5
Standard Softmax Transformer
HNC-SVD Compressed (50%)
Gated Linear Attention (GLA)
Gated DeltaNet (GDN)
SVD-Linear Attention (Ours)
Standard Softmax Transformer
SVD-Linear Attention (Ours)
Gated DeltaNet (GDN)
HNC-SVD Compressed (50%)
Gated Linear Attention (GLA)