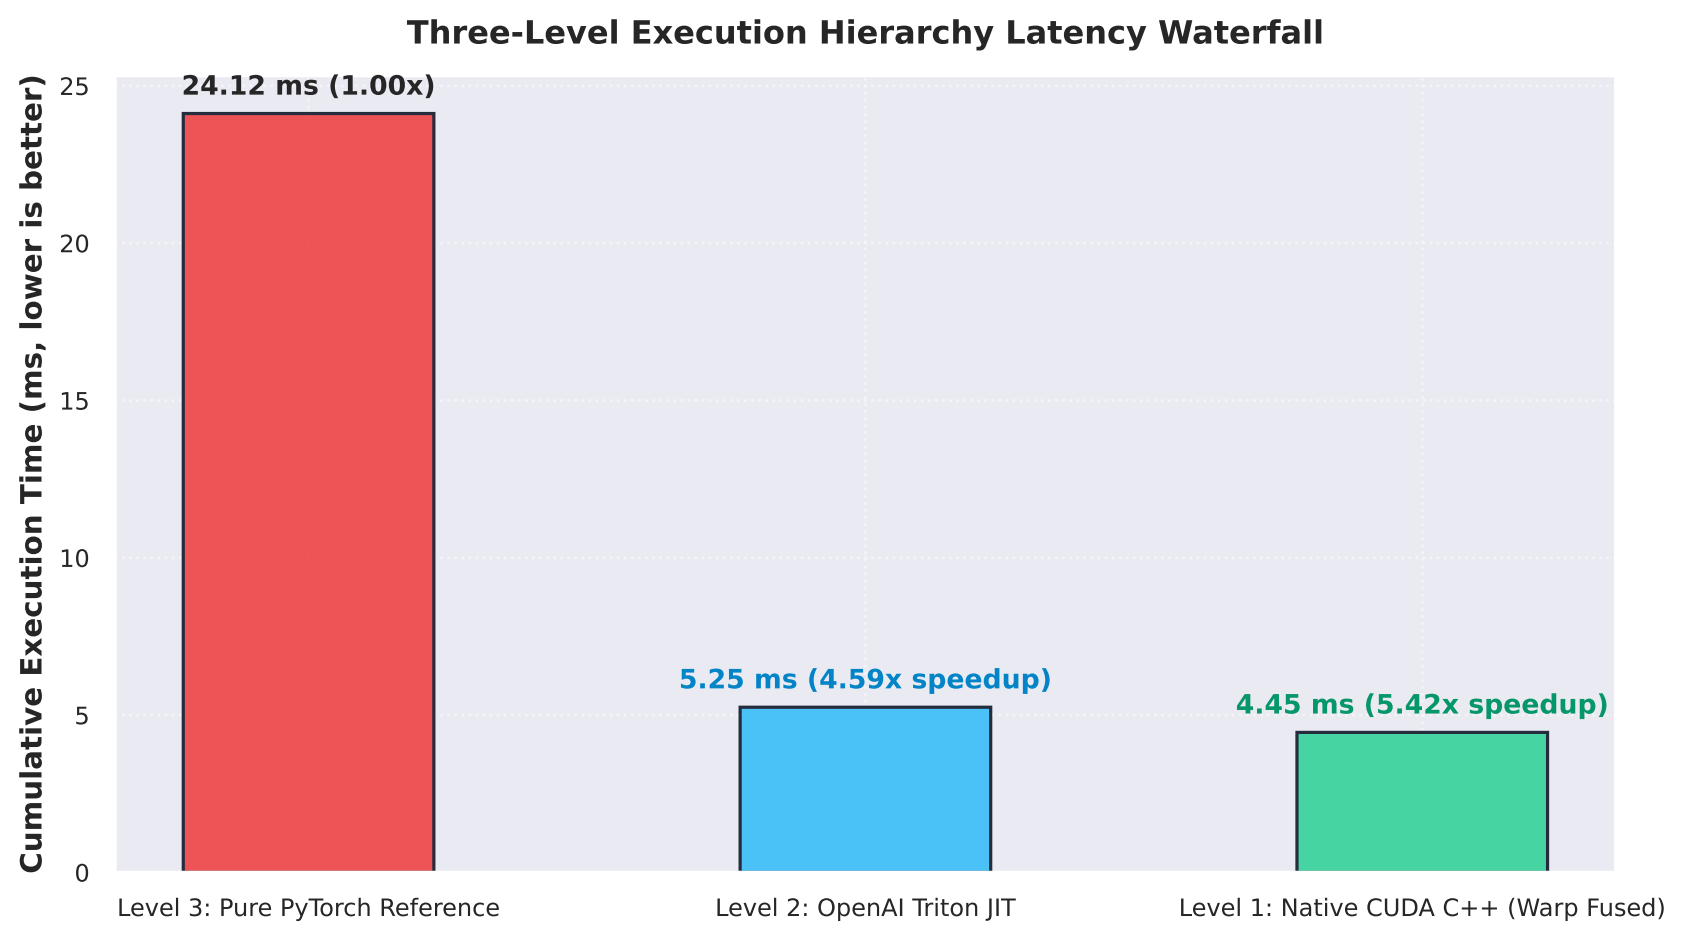
Three-Level Execution Hierarchy Latency Waterfall
Cumulative Execution Time (ms, lower is better)
0
5
10
15
20
25
24.12 ms (1.00x)
5.25 ms (4.59x speedup)
4.45 ms (5.42x speedup)
Level 3: Pure PyTorch Reference
Level 2: OpenAI Triton JIT
Level 1: Native CUDA C++ (Warp Fused)